

Exploring a data frame

Table of contents

1	readr, dplyr, and ggplot2	1
2	Read the Gapminder labor cost data set	2
3	dplyr::glimpse	2
4	summary	3
5	summary with categorical columns as factors	3
6	Rename a column with base R	4
7	Rename a column with dplyr	4
8	Plot the data set	4

1 readr, dplyr, and ggplot2

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

```
intersect, setdiff, setequal, union
```

```
library(readr)
library(ggplot2)
```

These are the three libraries you need most when you explore a tabular dataset.

2 Read the Gapminder labor cost data set

```
myfilepath <- |
  ↪ "datasets_ATRIUM/gapminder_hourly_labour_cost_constant_2017_usd--by--geo--time.csv"
laborcost_df <- read_csv(file = myfilepath,
                        show_col_types = TRUE)
```

Rows: 548 Columns: 3

-- Column specification -----

Delimiter: ","

chr (1): geo

dbl (2): time, hourly_labour_cost_constant_2017_usd

i Use `spec()` to retrieve the full column specification for this data.

i Specify the column types or set `show_col_types = FALSE` to quiet this message.

It is the file demonstrated in the previous session. We saved it into the folder `datasets_ATRIUM`. You can also get it at https://raw.githubusercontent.com/open-numbers/ddf-gapminder-systema-globalis/refs/heads/master/countries-etc-datapoints/ddf-datapoints-hourly_labour_cost_constant_2017_usd-by-geo-time.csv.

Watch the message `read_csv` gives you about the file. You can suppress it by overriding the default to `show_col_types = FALSE`.

3 `dplyr::glimpse`

- peek at the dataset (tilted 90°)

```
glimpse(laborcost_df)
```

```

Rows: 548
Columns: 3
$ geo                                <chr> "arg", "arg", "arm", "arm", "arm"~
$ time                              <dbl> 2011, 2012, 2011, 2012, 2013, 201~
$ hourly_labour_cost_constant_2017_usd <dbl> 0.92, 1.04, 4.23, 4.59, 6.12, 6.0~

```

Gives you the number of rows and columns, the column names with their data class, and it also shows as many elements (values) in each column as to fit your screen.

4 summary

```
summary(laborcost_df)
```

geo	time	hourly_labour_cost_constant_2017_usd
Length:548	Min. :1994	Min. : 0.000
Class :character	1st Qu.:2005	1st Qu.: 9.867
Mode :character	Median :2011	Median :18.320
	Mean :2010	Mean :19.686
	3rd Qu.:2017	3rd Qu.:26.915
	Max. :2020	Max. :48.720

Gives you the “five-number summary” of each numeric column (it’s often called this way, although the numbers are obviously six...). With categorical columns, it depends, whether the column is a character vector or a factor.

5 summary with categorical columns as factors

geo	time	hourly_labour_cost_constant_2017_usd
cze : 22	Min. :1994	Min. : 0.000
svn : 22	1st Qu.:2005	1st Qu.: 9.867
cyp : 21	Median :2011	Median :18.320
deu : 21	Mean :2010	Mean :19.686
pol : 21	3rd Qu.:2017	3rd Qu.:26.915
svk : 21	Max. :2020	Max. :48.720
(Other):420		

If you have a data frame with categorical variables converted to factors, the summary will show you a glimpse of their **levels** (unique values) and their frequencies, as well as tell you how many levels there are.

So far, do not worry about factors. The `dplyr` as well as the `ggplot2` libraries do this factor conversion on the fly whenever they need it.

6 Rename a column with base R

hourly_labour_cost_constant_2017_usd too long, shorten to labor_cost.

```
colnames(laborcost_df)[colnames(laborcost_df) ==  
                        "hourly_labour_cost_constant_2017_usd"] <-  
  ↪ "labor_cost"
```

```
colnames(laborcost_df)
```

```
[1] "geo"      "time"     "labor_cost"
```

7 Rename a column with dplyr

```
laborcost_df <- rename(.data = laborcost_df,  
                       labor_cost =  
  ↪ hourly_labour_cost_constant_2017_usd  
                       )  
colnames(laborcost_df)
```

```
[1] "geo"      "time"     "labor_cost"
```

You already know you could have named all columns your way when reading in the file. Here are two ways to rename a column: one base-R-like and the other one provided by dplyr.

8 Plot the data set

```
ggplot(data = laborcost_df,  
       mapping = aes(x = time,  
                     y = labor_cost,  
                     color = geo)) + geom_point()
```

