

tidyr and stringr

Silvie Cinková

2025-08-13

Table of contents

1	Libraries	1
2	Data	2
3	tidyr	2
4	tidyr::unite (before)	3
5	tidyr::unite (after)	3
6	tidyr::separate	4
7	tidyr::pivot_wider	5
8	tidyr::pivot_longer	6
9	Separating into rows	7
10	continuation	7

1 Libraries

```
library(readr)
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
library(tidyr)
library(magrittr)
```

Attaching package: 'magrittr'

The following object is masked from 'package:tidyr':

extract

```
library(gapminder)
library(ggplot2)
library(tidytext)
```

2 Data

```
billionaires_df <- read_tsv("datasets_ATRIUM/billionaires_combined.tsv",
  ↪ show_col_types = FALSE)
bil_unite <- billionaires_df %>% distinct(name.x, birth_comb) %>%
  ↪ drop_na(birth_comb)
bil_separate <- billionaires_df %>% distinct(name.x, person) %>%
  ↪ drop_na(name.x)
gap_cze_ger_gdp <- gapminder %>%
  filter(country %in% c("Czech Republic", "Germany")) %>%
  select(country, year, gdpPercap)
```

3 tidyr

- lumps and splits column values into new columns

- transforms several variable columns into categories of a new variable and the other way round
- completes observations with **missing values**
- manages **nested columns** (vector/list inside column!)

This library is a “heavy-duty” assistant to `dplyr`. You use it when you need to make your data tidy for your purposes. Remember, tidy means that each row represents an observation and each column represents one variable. Sometimes it is a matter of perspective. Maybe we could say that the data is tidy when you can map all relevant variable to ggplot aesthetic scales.

The original four big verbs of tidyr were two pairs of twins: `unite` and `separate`, and `pivot_longer` (formerly `gather`) and `pivot_wider` (formerly `spread`)

4 tidyr::unite (before)

```
bil_unite %>% slice(1:10)
```

```
# A tibble: 10 x 2
  name.x                birth_comb
  <chr>                <dbl>
1 A. Jerrold Perenchio    1931
2 Abdulla bin Ahmad Al Ghurair    1955
3 Abdullah bin Sulaiman Al Rajhi    1929
4 Abdulsamad Rabiou    1960
5 Abhay Soi    1973
6 Abhay Vakil    1952
7 Abigail Johnson    1962
8 Abilio dos Santos Diniz    1937
9 Achal Bakeri    1961
10 Acharya Balakrishna    1972
```

5 tidyr::unite (after)

```

bil_unite %>% slice(1:10) %>%
  unite(col = ID, name.x, birth_comb,
    ↪ remove = FALSE) %>%
  kableExtra::kable()
sep = "***",

```

ID	name.x	birth_comb
A. Jerrold Perenchio***1931	A. Jerrold Perenchio	1931
Abdulla bin Ahmad Al Ghurair***1955	Abdulla bin Ahmad Al Ghurair	1955
Abdullah bin Sulaiman Al Rajhi***1929	Abdullah bin Sulaiman Al Rajhi	1929
Abdulsamad Rabiou***1960	Abdulsamad Rabiou	1960
Abhay Soi***1973	Abhay Soi	1973
Abhay Vakil***1952	Abhay Vakil	1952
Abigail Johnson***1962	Abigail Johnson	1962
Abilio dos Santos Diniz***1937	Abilio dos Santos Diniz	1937
Achal Bakeri***1961	Achal Bakeri	1961
Acharya Balakrishna***1972	Acharya Balakrishna	1972

6 tidyr::separate

- splits values into two or more new columns

```

bil_separate %>% slice(1:2) %>%
  separate(col = name.x, into = c("firstnames", "middlenames", "lastnames" ),
    ↪ sep = " ", fill = "left")

```

```

# A tibble: 2 x 4
  firstnames middlenames lastnames person
  <chr>      <chr>      <chr>    <chr>
1 A.        Jayson      Adair    a_jayson_adair
2 A.        Jerrold     Perenchio a_jerrold_perenchio

```

```

set.seed(15)
bil_separate %>% slice_sample(n = 10) %>%
  separate(col = name.x,
    into = c("firstnames", "lastnames" ),
    sep = " ", fill = "left", remove = FALSE)

```

```
# A tibble: 10 x 4
```

	name.x	firstnames	lastnames	person
	<chr>	<chr>	<chr>	<chr>
1	Xu Xiaoqun	Xu	Xiaoqun	xu_xiaoqun
2	Irwin Jacobs	Irwin	Jacobs	irwin_jacobs
3	Wang Chaobin	Wang	Chaobin	wang_chaobin
4	Anton Schlecker	Anton	Schlecker	anton_schlecker
5	Thomas James	Thomas	James	thomas_james
6	GSK Velu	GSK	Velu	gsk_velu
7	Charoen Sirivadhanabhakdi	Charoen	Sirivadhanabhakdi	charoen_sirivadhanabh~
8	Yuriy Kosiuk	Yuriy	Kosiuk	yuriy_kosiuk
9	Elaine Marshall	Elaine	Marshall	elaine_marshall
10	Jaime Botin	Jaime	Botin	jaime_botin_1

In some functions in `tidyr`, but also in `dplyr`, you will benefit of a good command of string operations. That is, when you can find general patterns in strings, mainly with Regular Expressions. When your data is structured very well, it can be easy; such as separate first names from surnames, when they are all in the same order and separated by comma. Most often this is not the case, but you can at least approximate a good result with a more advanced pattern. The current version of `tidyr` even has a separate function that works with regular expressions, but I was not able to get with it smarter results than with the ordinary one that only works with the delimiter.

At any rate, R has a wonderful library for work with strings: `stringr`. I will include `stringr` functions wherever relevant.

7 tidyr::pivot_wider

```
gap_czger_gdp_wide <- gap_cze_ger_gdp %>%  
  pivot_wider(names_from = country, values_from = gdpPercap)  
gap_czger_gdp_wide %>% slice(1:3) %>% kableExtra::kable()
```

year	Czech Republic	Germany
1952	6876.140	7144.114
1957	8256.344	10187.827
1962	10136.867	12902.463

```
gap_czger_gdp_wide %$%  
  cor(x = `Czech Republic`, y = Germany, method = "pearson")
```

```
[1] 0.946459
```

Imagine you want to compute a correlation of the temporal development between two countries, considering an indicator, such as GDP per capita. To compute the correlation, you must have the data in two separate variables. This is how you would prepare `gapminder` to correlate GDP per capita 1952 - 2007 between Czechia and Germany.

Format changes in data frames: wider = fewer rows and more columns, longer = (sometimes) fewer columns but definitely more rows, both compared to the state before the manipulation.

8 `tidyr::pivot_longer`

```
gap_czger_gdp_wide %>% slice(1:10) %>%  
  pivot_longer(cols = c(`Czech Republic`, Germany), names_to = "COUNTRY",  
    ↪ values_to = "GDPperCap")
```

```
# A tibble: 20 x 3
```

	year	COUNTRY	GDPperCap
	<int>	<chr>	<dbl>
1	1952	Czech Republic	6876.
2	1952	Germany	7144.
3	1957	Czech Republic	8256.
4	1957	Germany	10188.
5	1962	Czech Republic	10137.
6	1962	Germany	12902.
7	1967	Czech Republic	11399.
8	1967	Germany	14746.
9	1972	Czech Republic	13108.
10	1972	Germany	18016.
11	1977	Czech Republic	14800.
12	1977	Germany	20513.
13	1982	Czech Republic	15377.
14	1982	Germany	22032.
15	1987	Czech Republic	16310.
16	1987	Germany	24639.
17	1992	Czech Republic	14297.
18	1992	Germany	26505.
19	1997	Czech Republic	16049.
20	1997	Germany	27789.

... but you definitely prefer the GDP in one column and countries in the other column when you want to plot the development comparison. Wider tables are often considered to be more human-readable, and therefore you typically get them from sources that were primarily designed for print.

9 Separating into rows

```
industry_terms <- billionaires_df %>%
  select(c(person, time, countries, industry, income_groups, world_6region))
  ↪ %>%
  separate_longer_delim(cols = c("industry"), delim = ";")
set.seed(33)
industry_terms %>% slice_sample(n = 10)
```

A tibble: 10 x 6

	person <chr>	time <dbl>	countries <chr>	industry <chr>	income_groups <chr>	world_6region <chr>
1	yang_huiyan	2022	chn	"Real Estate, Edu~	upper_middle~	east_asia_pa~
2	horst_wortmann	2016	deu	"Apparel & Textil~	high_income	europa_centra~
3	wang_wenxue	2018	chn	"Real Estate"	upper_middle~	east_asia_pa~
4	b_wayne_hughes	2005	usa	"Service"	high_income	america
5	daniela_herz	2013	deu	"Investments"	high_income	europa_centra~
6	edward_lampert	2017	usa	" Finance"	high_income	america
7	xu_jingren	2020	chn	"Healthcare"	upper_middle~	east_asia_pa~
8	qiu_guanghe	2021	chn	"Apparel"	upper_middle~	east_asia_pa~
9	elon_musk	2019	usa	"Automotive"	high_income	america
10	ennio_doris	2009	ita	" Finance"	high_income	europa_centra~

10 continuation

```
industry_terms %<>%
  separate_longer_delim(cols = c("industry"), delim = "&") %>%
  separate_longer_delim(cols = c("industry"), delim = "and") %>%
  separate_longer_delim(cols = c("industry"), delim = ",")
set.seed(10)
industry_terms %>% slice_sample(n = 10)
```

A tibble: 10 x 6

	person	time	countries	industry	income_groups	world_6region
	<chr>	<dbl>	<chr>	<chr>	<chr>	<chr>
1	christy_walton	2013	usa	"Fashion "	high_income	america
2	ty_warner	2004	usa	"Hospital~	high_income	america
3	mustafa_rahmi_koc	2018	tur	" Gas"	upper_middle~	europe_cent~
4	oleg_deripaska	2008	rus	"Metals "	upper_middle~	europe_cent~
5	doris_fisher	2021	usa	" Fashion~	high_income	america
6	gerald_ford	2017	usa	"Finance "	high_income	america
7	gregorio_perez_companc	2005	arg	" Food"	upper_middle~	america
8	robert_bass	2012	usa	"Energy"	high_income	america
9	norma_lerner	2018	usa	"Financia~	high_income	america
10	yuri_kovalchuk	2018	rus	"Finance "	upper_middle~	europe_cent~

```
industry_tf_idf <- industry_terms %>% filter(nchar(industry) > 1) %>%
  group_by(industry, countries, world_6region) %>%
  count(name = "freq") %>% ungroup() %>%
  tidytext::bind_tf_idf(term = industry,
                        document = countries,
                        n = freq) %>%
  group_by(countries, world_6region) %>%
  slice_max(order_by = tf_idf, n = 3) %>% ungroup()
```

```
library(dplyr)
library(tidygraph)
```

Attaching package: 'tidygraph'

The following object is masked from 'package:stats':

filter

```
library(ggraph)

# Prepare edge list
edges <- industry_tf_idf %>%
  select(countries, industry, tf_idf)

# Create node list
nodes <- tibble(name = unique(c(edges$countries, edges$industry))) %>%
  mutate(type = if_else(name %in% edges$countries, "country", "industry"))

# Create graph object
```




```
graph <- tbl_graph(nodes = nodes, edges = edges, directed = FALSE)
graph
```

```
# A tbl_graph: 176 nodes and 204 edges
#
# A bipartite simple graph with 11 components
#
# Node Data: 176 x 2 (active)
  name  type
  <chr> <chr>
1 ago   country
2 are   country
3 arg   country
4 aus   country
5 aut   country
6 bel   country
7 bgd   country
8 bgr   country
9 bra   country
10 brb  country
# i 166 more rows
#
# Edge Data: 204 x 3
  from  to  tf_idf
  <int> <int> <dbl>
1     1   74 0.168
2     1   75 0.115
3     1   76 0.0968
# i 201 more rows
```

```
# Plot with ggraph
ggraph(graph, layout = "auto") +
  geom_edge_link(aes(edge_width = tf_idf), alpha = 0.6) +
  geom_node_point(aes(color = type), size = 0.3) +
  geom_node_text(aes(label = name, color = type), repel = TRUE, size = 2,
  ↪  max.overlaps = 100) +
  scale_edge_width(range = c(0.2, 2)) +
  theme_void() +
  labs(title = "Bipartite Network: Countries and Industries")
```

Using "stress" as default layout

Bipartite Network: Countries and Industries



```
ggsave(filename = "my_output_files/billionairs_industry_country_ggraph.pdf",
  width = 7 * 2, height = 7 * 2.2)
```

Warning in grid.Call.graphics(C_text, as.graphicsAnnot(x\$label), x\$x, x\$y, :
conversion failure on 'Paper Packaging ' in 'mbscsToSbcs': for (U+3001)